

11: Generalized Linear Models

Environmental Data Analytics / Kateri Salk

Spring 2019

LESSON OBJECTIVES

1. Describe the components of the generalized linear model (GLM)
2. Apply special cases of the GLM to real datasets
3. Interpret and report the results of GLMs in publication-style formats

SET UP YOUR DATA ANALYSIS SESSION

```
getwd()

## [1] "/Users/laurie/Desktop/Envtl_Data_Analytics/MuzzyGitFile"

library(tidyverse)

PeterPaul.nutrients <- read.csv("./Data/Processed/NTL-LTER_Lake_Nutrients_PeterPaul_Processed.csv")
EPAair <- read.csv("./Data/Processed/EPAair_03PM25_3sites1718_processed.csv")

# Set date to date format
EPAair$Date <- as.Date(EPAair$Date, format = "%Y-%m-%d")
PeterPaul.nutrients$sampldate <- as.Date(PeterPaul.nutrients$sampldate, format = "%Y-%m-%d")

# remove negative values for depth_id
PeterPaul.nutrients <- filter(PeterPaul.nutrients, depth_id > 0)

# set depth_id to factor
PeterPaul.nutrients$depth_id <- as.factor(PeterPaul.nutrients$depth_id) #setting it as factor
class(PeterPaul.nutrients$depth_id) #factor (used to be integer)

## [1] "factor"

mytheme <- theme_classic(base_size = 14) +
  theme(axis.text = element_text(color = "black"),
        legend.position = "top")
theme_set(mytheme)
```

GENERALIZED LINEAR MODELS

The one-sample test (model of the mean), two-sample t-test, analysis of variance (ANOVA), and linear regression are all special cases of the **generalized linear model** (GLM). The GLM also includes analyses not covered in this class, including logistic regression, multinomial regression, chi square, and log-linear models. The common characteristic of general linear models is the expression of a continuous response variable as a linear combination of the effects of categorical or continuous explanatory variables, plus an error term that expresses the random error associated with the coefficients of all explanatory variables. The explanatory variables comprise the deterministic component of the model, and the error term comprises the stochastic component of the model. Historically, artificial distinctions were made between linear models that contained categorical and continuous explanatory variables, but this distinction is no longer made. The inclusion of

these models within the umbrella of the GLM allows models to fit the main effects of both categorical and continuous explanatory variables as well as their interactions.

Choosing a model from your data: A “cheat sheet”

T-test: Continuous response, one categorical explanatory variable with two categories (or comparison to a single value if a one-sample test).

One-way ANOVA (Analysis of Variance): Continuous response, one categorical explanatory variable with more than two categories.

Two-way ANOVA (Analysis of Variance) Continuous response, two categorical explanatory variables.

Single Linear Regression Continuous response, one continuous explanatory variable.

Multiple Linear Regression Continuous response, two or more continuous explanatory variables.

ANCOVA (Analysis of Covariance) Continuous response, categorical explanatory variable(s) and continuous explanatory variable(s).

If multiple explanatory variables are chosen, they may be analyzed with respect to their **main effects** on the model (i.e., their separate impacts on the variance explained) or with respect to their **interaction effects**, the effect of interacting explanatory variables on the model.

Assumptions of the GLM

The GLM is based on the assumption that the data approximate a normal distribution (or a linearly transformed normal distribution). We will discuss the non-parametric analogues to several of these tests if the assumptions of normality are violated. For tests that analyze categorical explanatory variables, the assumption is that the variance in the response variable is equal among groups. Note: environmental data often violate the assumptions of normality and equal variance, and we will often proceed with a GLM even if these assumptions are violated. In this situation, you must justify your decision.

T-TEST AND ONE-WAY ANOVA

One-sample t-test

The object of a one sample test is to test the null hypothesis that the mean of the group is equal to a specific value. For example, we might ask ourselves (from the EPA air quality processed dataset):

Are Ozone levels below the threshold for “good” AQI index (0-50)?

```
summary(EPAair$Ozone)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.     NA's
##      5.00   31.00   37.00   36.92   44.00   97.00     868
```

```
# Evaluate assumption of normal distribution
```

```
shapiro.test(EPAair$Ozone) #will show us the distr, and whether or not violating assumptions of normality
```

```
##
```

```
## Shapiro-Wilk normality test
```

```
##
```


```
## data: EPAair$Ozone
```

```
## W = 0.97317, p-value = 2.747e-13
```

```
ggplot(EPAair, aes(x = Ozone)) +  
  geom_histogram(stat = "count")
```


```
## Warning: Ignoring unknown parameters: binwidth, bins, pad
```

```
## Warning: Removed 868 rows containing non-finite values (stat_count).
```



```
qqnorm(EPAair$Ozone); qqline(EPAair$Ozone) #qqplot shows that when points fall off the line, they're fu
```

Normal Q-Q Plot



#v low p-value, histogram shows that it's not really normal, but not bimodal

```
O3.onesample <- t.test(EPAair$Ozone, mu = 50, alternative = "less") # mu is for what we're comparing it
O3.onesample #making it an r object so we can call it up again
```

```
##
## One Sample t-test
##
## data: EPAair$Ozone
## t = -41.911, df = 1084, p-value < 2.2e-16
## alternative hypothesis: true mean is less than 50
## 95 percent confidence interval:
##      -Inf 37.43006
## sample estimates:
## mean of x
## 36.91613
```

What information does the output give us? How might we report this information in a report?

ANSWER: We have a really low p-value, so we can reject the null hyp. The t-score is also low, so # we have a 95% confidence that the mean Ozone is btwn -infinity and 37 ug/L #look at t-stat and degrees of freedom We'd report that the Ozone AQI values in 2017-18 were signif lower than 50 (1-sample t-test, $t = -41.9$ and $p < 0.0001$, $df = 1084$), in the "good" threshold.

Two-sample t-test

The two-sample t test is used to test the hypothesis that the mean of two samples is equivalent. Unlike the one-sample tests, a two-sample test requires a second assumption that the variance of the two groups is equivalent. Are Ozone levels different between Blackstone and Bryson City?

```

shapiro.test(EPAair$Ozone[EPAair$Site.Name == "Blackstone"])

##
## Shapiro-Wilk normality test
##
## data: EPAair$Ozone[EPAair$Site.Name == "Blackstone"]
## W = 0.97221, p-value = 6.349e-09

shapiro.test(EPAair$Ozone[EPAair$Site.Name == "Bryson City"])

##
## Shapiro-Wilk normality test
##
## data: EPAair$Ozone[EPAair$Site.Name == "Bryson City"]
## W = 0.97189, p-value = 2.228e-08

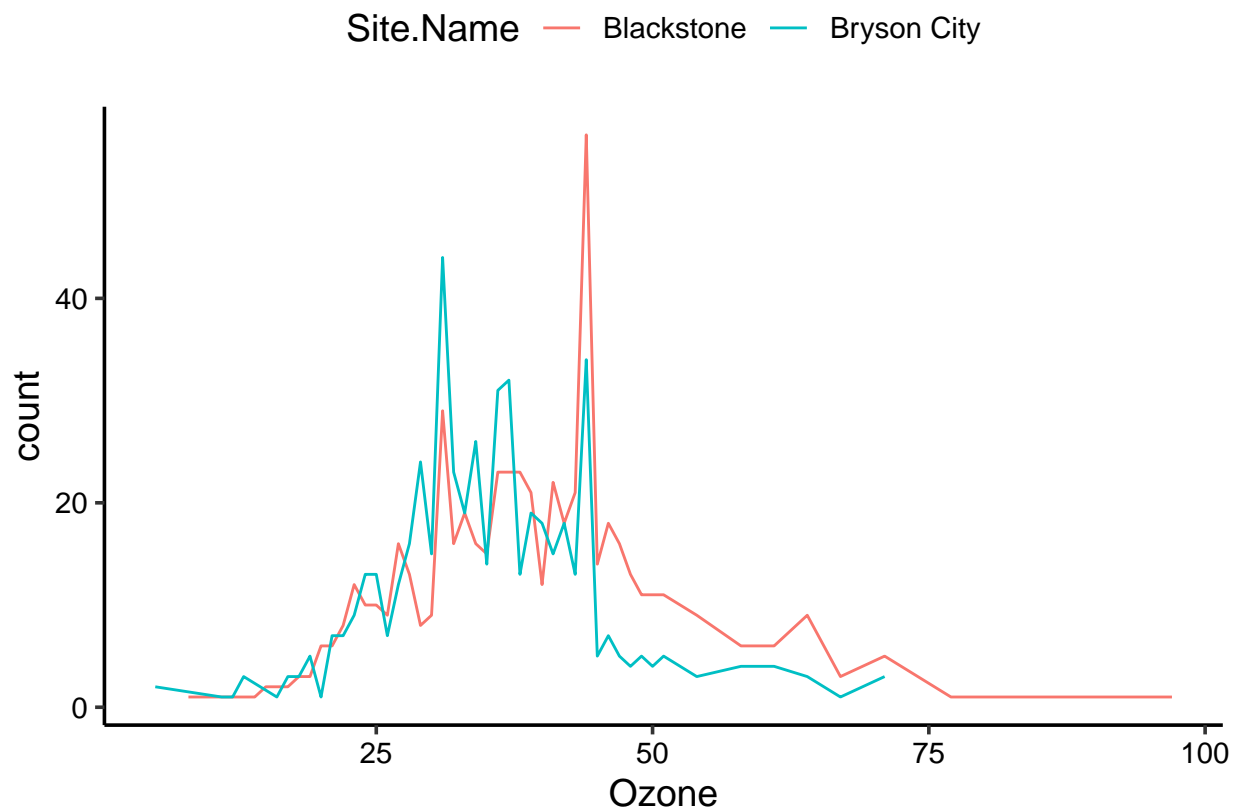
var.test(EPAair$Ozone ~ EPAair$Site.Name) #gives us f-stat and p-value

##
## F test to compare two variances
##
## data: EPAair$Ozone by EPAair$Site.Name
## F = 1.3678, num df = 569, denom df = 514, p-value = 0.0002955
## alternative hypothesis: true ratio of variances is not equal to 1
## 95 percent confidence interval:
##  1.154854 1.618780
## sample estimates:
## ratio of variances
##      1.367782
#the variance btwn these 2 populations is different

ggplot(EPAair, aes(x = Ozone, color = Site.Name)) +
  geom_freqpoly(stat = "count")

## Warning: Removed 868 rows containing non-finite values (stat_count).

```



```
# Format as a t-test
O3.twosample <- t.test(EPAair$Ozone ~ EPAair$Site.Name)
O3.twosample #we can reject the null that ; the p-val is really high
```

```
##
## Welch Two Sample t-test
##
## data: EPAair$Ozone by EPAair$Site.Name
## t = 5.3875, df = 1079.8, p-value = 8.766e-08
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  2.098082 4.501782
## sample estimates:
## mean in group Blackstone mean in group Bryson City
##           38.48246           35.18252
```

```
O3.twosample$p.value #use dollar sign to call up any portion of the test, so we can generate tables to
```

```
## [1] 8.765983e-08
```

```
# Format as a GLM
O3.twosample2 <- lm(EPAair$Ozone ~ EPAair$Site.Name)
summary(O3.twosample2) #not all info is superuseful but could be handy
```

```
##
## Call:
## lm(formula = EPAair$Ozone ~ EPAair$Site.Name)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -30.482  -6.183  -0.183   5.518  58.518
##
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)      38.4825     0.4253  90.477 < 2e-16 ***
## EPAair$Site.NameBryson City  -3.2999     0.6174  -5.345  1.1e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 10.15 on 1083 degrees of freedom
## (868 observations deleted due to missingness)
## Multiple R-squared:  0.0257, Adjusted R-squared:  0.0248
## F-statistic: 28.57 on 1 and 1083 DF, p-value: 1.101e-07
```

Non-parametric equivalent of t-test: Wilcoxon test

When we wish to avoid the assumption of normality, we can apply *distribution-free*, or non-parametric, methods in the form of the Wilcoxon rank sum (Mann-Whitney) test. The Wilcoxon test replaces the data by their rank and calculates the sum of the ranks for each group. Notice that the output of the Wilcoxon test is more limited than its parametric equivalent.

```
03.onesample.wilcox <- wilcox.test(EPAair$Ozone, mu = 50, alternative = "less")
03.onesample.wilcox #much more limited, but could help if we've really diverged from assumptions of norm
```

```
##
## Wilcoxon signed rank test with continuity correction
##
## data: EPAair$Ozone
## V = 25828, p-value < 2.2e-16
## alternative hypothesis: true location is less than 50
03.twosample.wilcox <- wilcox.test(EPAair$Ozone ~ EPAair$Site.Name)
03.twosample.wilcox
```

```
##
## Wilcoxon rank sum test with continuity correction
##
## data: EPAair$Ozone by EPAair$Site.Name
## W = 175960, p-value = 1.451e-08
## alternative hypothesis: true location shift is not equal to 0
```

One-way ANOVA

A one-way ANOVA is the same test in practice as a two-sample t-test but for three or more groups. In R, we can run the model with the function `lm` or `aov`, the latter of which will allow us to run post-hoc tests to determine pairwise differences.

Are PM2.5 levels different between Blackstone, Bryson City, and Triple Oak?

```
shapiro.test(EPAair$PM2.5[EPAair$Site.Name == "Blackstone"])
```

```
##
## Shapiro-Wilk normality test
##
## data: EPAair$PM2.5[EPAair$Site.Name == "Blackstone"]
```

```
## W = 0.99335, p-value = 0.01489
```

```
shapiro.test(EPAair$PM2.5[EPAair$Site.Name == "Bryson City"])
```

```
##
```

```
## Shapiro-Wilk normality test
```

```
##
```

```
## data: EPAair$PM2.5[EPAair$Site.Name == "Bryson City"]
```

```
## W = 0.98207, p-value = 2.527e-07
```

```
shapiro.test(EPAair$PM2.5[EPAair$Site.Name == "Triple Oak"])
```

```
##
```

```
## Shapiro-Wilk normality test
```

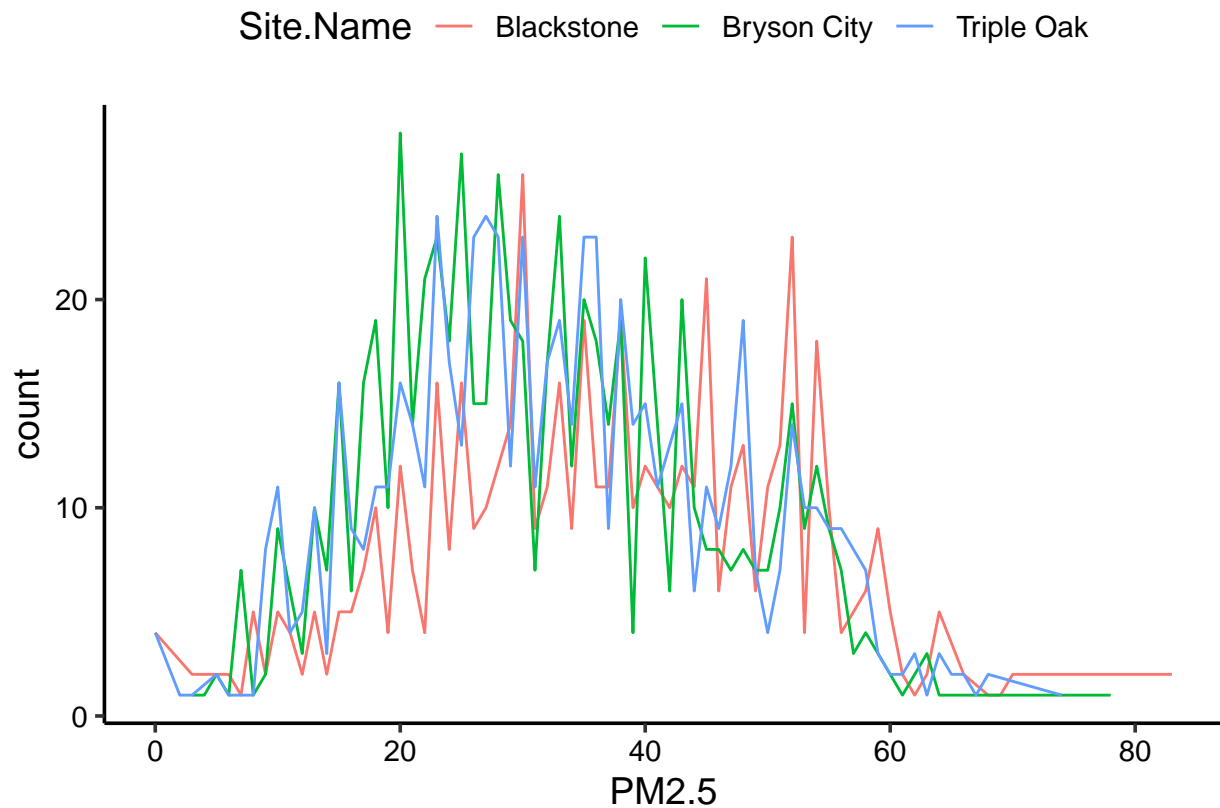
```
##
```

```
## data: EPAair$PM2.5[EPAair$Site.Name == "Triple Oak"]
```

```
## W = 0.99064, p-value = 0.0002744
```


```
ggplot(EPAair, aes(x = PM2.5, color = Site.Name)) +  
  geom_freqpoly(stat = "count") #no crazy tails, so maybe we should proceed w /ANOVA
```

```
## Warning: Removed 52 rows containing non-finite values (stat_count).
```



```
qqnorm(EPAair$PM2.5); qqline(EPAair$PM2.5) #the variability is around the edges (common for env data)
```


Normal Q-Q Plot



```
bartlett.test(EPAair$PM2.5 ~ EPAair$Site.Name) #k-sq value and leaves 2 df; this test shows that these
```

```
##
## Bartlett test of homogeneity of variances
##
## data: EPAair$PM2.5 by EPAair$Site.Name
## Bartlett's K-squared = 4.9951, df = 2, p-value = 0.08229
# Format as a GLM
PM2.5.anova <- lm(EPAair$PM2.5 ~ EPAair$Site.Name) # ~ explanatory variable
summary(PM2.5.anova)
```

```
##
## Call:
## lm(formula = EPAair$PM2.5 ~ EPAair$Site.Name)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-36.726	-10.300	-0.726	10.274	46.274

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	36.7261	0.5902	62.231	< 2e-16 ***
EPAair\$Site.NameBryson City	-4.4266	0.7977	-5.549	3.28e-08 ***
EPAair\$Site.NameTriple Oak	-3.2461	0.7967	-4.075	4.80e-05 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 13.9 on 1898 degrees of freedom
## (52 observations deleted due to missingness)
```


```
## Multiple R-squared:  0.01674,    Adjusted R-squared:  0.01571
## F-statistic: 16.16 on 2 and 1898 DF,  p-value: 1.1e-07
# Format as an aov
PM2.5.anova2 <- aov(EPAair$PM2.5 ~ EPAair$Site.Name)
summary(PM2.5.anova2)

##              Df Sum Sq Mean Sq F value   Pr(>F)
## EPAair$Site.Name    2    6247   3123.6    16.16 1.1e-07 ***
## Residuals          1898  366884    193.3
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 52 observations deleted due to missingness
# Run a post-hoc test for pairwise differences
TukeyHSD(PM2.5.anova2) #tells us how different the sites are from each other

##      Tukey multiple comparisons of means
##      95% family-wise confidence level
##
## Fit: aov(formula = EPAair$PM2.5 ~ EPAair$Site.Name)
##
## $`EPAair$Site.Name`
##              diff              lwr              upr              p adj
## Bryson City-Blackstone -4.426573 -6.2976740 -2.555472 0.0000001
## Triple Oak-Blackstone  -3.246126 -5.1147155 -1.377537 0.0001419
## Triple Oak-Bryson City   1.180447 -0.5972964   2.958191 0.2645306

# Plot the results
# How might you edit this graph to make it attractive?
# How might you illustrate significant differences?
PM2.5.anova.plot <- ggplot(EPAair, aes(x = Site.Name, y = PM2.5)) +
  geom_violin(draw_quantiles = 0.5)
print(PM2.5.anova.plot)

## Warning: Removed 52 rows containing non-finite values (stat_ydensity).
```



What information does the output give us? How might we report this information in a report?

ANSWER: we'd want to show stat dif by maybe an asterisk by Blackston's mean, or put the numbers above the violins (or use `geom_text`). We did an ANOVA test, p-val was <0.0001 , $F=16.16$, $df = 1898$. Over that study period, BC and TO did not have stat significant diff

Non-parametric equivalent of ANOVA: Kruskal-Wallis Test

As with the Wilcoxon test, the Kruskal-Wallis test is the non-parametric counterpart to the one-way ANOVA. Here, the data from two or more independent samples are replaced with their ranks without regard to the grouping AND based on the between-group sum of squares calculations.

For multiple comparisons, a p-value < 0.05 indicates that there is a significant difference between groups, but it does not indicate which groups, or in this case, months, differ from each other.

To analyze specific pairs in the data, you must use a *post hoc* test. These include the Dunn's test, a pairwise Mann-Whitney with the Bonferroni correction, or the Conover-Iman test.

```
PM2.5.kw <- kruskal.test(EPAair$PM2.5 ~ EPAair$Site.Name)
PM2.5.kw #not getting linear bc we're not using parametric approach
```

```
##
## Kruskal-Wallis rank sum test
##
## data: EPAair$PM2.5 by EPAair$Site.Name
## Kruskal-Wallis chi-squared = 34.737, df = 2, p-value = 2.864e-08
# There are two functions to run the Dunn Test
# dunn.test(EPAair$PM2.5, EPAair$Site.Name, kw = T,
```

```
#           table = F, list = T, method = "holm", altp = T) #From package dunn.test
# dunnTest(EPAair$PM2.5, EPAair$Site.Name) #From package FSA
```

TWO-WAY ANOVA

Main effects

A two-way ANOVA allows us to examine the effects of two categorical explanatory variables on a continuous response variable. Let's look at the NTL-LTER nutrient dataset for Peter and Paul lakes. What if we wanted to know if total nitrogen concentrations differed based on lake and depth?

```
TNanova.main <- lm(PeterPaul.nutrients$tn_ug ~ PeterPaul.nutrients$lakename + PeterPaul.nutrients$depth_id,
summary(TNanova.main) #predicts 309ug/L for Paul Lk (the intercept) and 105 for Peter
```

```
##
## Call:
## lm(formula = PeterPaul.nutrients$tn_ug ~ PeterPaul.nutrients$lakename +
##     PeterPaul.nutrients$depth_id)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -894.80  -98.28  -37.18   60.55 2223.54
##
## Coefficients:
##                                Estimate Std. Error t value
## (Intercept)                   309.39      12.48  24.786
## PeterPaul.nutrients$lakenamePeter Lake    105.29      13.89   7.580
## PeterPaul.nutrients$depth_id2              97.28      25.63   3.796
## PeterPaul.nutrients$depth_id3             113.40      25.54   4.440
## PeterPaul.nutrients$depth_id4              78.97      24.90   3.172
## PeterPaul.nutrients$depth_id5              22.47      26.25   0.856
## PeterPaul.nutrients$depth_id6              39.00      29.50   1.322
## PeterPaul.nutrients$depth_id7             859.48      21.52  39.931
##                                Pr(>|t|)
## (Intercept)                   < 2e-16 ***
## PeterPaul.nutrients$lakenamePeter Lake 6.20e-14 ***
## PeterPaul.nutrients$depth_id2          0.000153 ***
## PeterPaul.nutrients$depth_id3          9.71e-06 ***
## PeterPaul.nutrients$depth_id4          0.001546 **
## PeterPaul.nutrients$depth_id5          0.392172
## PeterPaul.nutrients$depth_id6          0.186319
## PeterPaul.nutrients$depth_id7          < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 262 on 1415 degrees of freedom
## (922 observations deleted due to missingness)
## Multiple R-squared:  0.5522, Adjusted R-squared:  0.5499
## F-statistic: 249.2 on 7 and 1415 DF,  p-value: < 2.2e-16
```

```
TNanova.main2 <- aov(PeterPaul.nutrients$tn_ug ~ PeterPaul.nutrients$lakename + PeterPaul.nutrients$depth_id,
summary(TNanova.main2)
```

```
##                                Df      Sum Sq Mean Sq F value    Pr(>F)
```


```
## PeterPaul.nutrients$lakename      1  4034942  4034942    58.8 3.23e-14 ***
## PeterPaul.nutrients$depth_id      6 115687621 19281270   281.0 < 2e-16 ***
## Residuals                        1415  97103398   68624
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 922 observations deleted due to missingness
```

```
TukeyHSD(TNanova.main2)
```

```
## Tukey multiple comparisons of means
## 95% family-wise confidence level
##
## Fit: aov(formula = PeterPaul.nutrients$tn_ug ~ PeterPaul.nutrients$lakename + PeterPaul.nutrients$depth_id)
##
## $`PeterPaul.nutrients$lakename`
##              diff          lwr          upr p adj
## Peter Lake-Paul Lake 106.4994 79.25437 133.7444 0
##
## $`PeterPaul.nutrients$depth_id`
##              diff          lwr          upr          p adj
## 2-1  97.28178    21.617077 172.94648 0.0029119
## 3-1 113.40580    37.992518 188.81908 0.0001959
## 4-1  78.98288     5.473012 152.49274 0.0258461
## 5-1  22.46140   -55.056737  99.97953 0.9788037
## 6-1  39.00303   -48.096701 126.10275 0.8416669
## 7-1 859.47649   795.924201 923.02879 0.0000000
## 3-2  16.12402   -81.518085 113.76613 0.9990113
## 4-2 -18.29890  -114.478514  77.88071 0.9977987
## 5-2 -74.82038  -174.097160  24.45640 0.2824951
## 6-2 -58.27875  -165.204802  48.64730 0.6763937
## 7-2 762.19472   673.393186 850.99625 0.0000000
## 4-3 -34.42292  -130.404869  61.55903 0.9397834
## 5-3 -90.94440  -190.029693   8.14089 0.0964544
## 6-3 -74.40277  -181.151057  32.34551 0.3786337
## 7-3 746.07070   657.483293 834.65810 0.0000000
## 5-4 -56.52148  -154.165899  41.12294 0.6100323
## 6-4 -39.97985  -145.392060  65.43236 0.9221509
## 7-4 780.49362   693.520832 867.46640 0.0000000
## 6-5  16.54163   -91.703900 124.78716 0.9993654
## 7-5 837.01510   746.629113 927.40108 0.0000000
## 7-6 820.47347   721.746941 919.20000 0.0000000
```

```
# Plot the results
# How might you edit this graph to make it attractive?
# How might you illustrate significant differences?
# imp't to visualize data: lots of N at bottom, means for Peter lk higher than Paul
TNanova.plot <- ggplot(PeterPaul.nutrients, aes(x = lakename, y = tn_ug, color = depth_id)) +
  geom_boxplot()
print(TNanova.plot)
```

```
## Warning: Removed 922 rows containing non-finite values (stat_boxplot).
```



Interaction effects

We may expect the effects of lake and depth to be dependent on each other. For instance, since `depth_id` is standardized across lakes, the concentrations at each `depth_id` might depend on which lake is sampled. In this case, we might choose to run an interaction effects two-way ANOVA, which will examine the individual effects of the explanatory variables as well as the interaction of the explanatory variables.

The output gives test statistics for each explanatory variable as well as the interaction effect of the explanatory variables. If the p-value for the interaction effect is less than 0.05, then we would consider the interaction among the explanatory variables to be significant.

```
TNanova.interaction <- aov(PeterPaul.nutrients$tn_ug ~ PeterPaul.nutrients$lakename * PeterPaul.nutrients$depth_id)
summary(TNanova.interaction)
```

```
##                                Df    Sum Sq
## PeterPaul.nutrients$lakename      1  4034942
## PeterPaul.nutrients$depth_id       6 115687621
## PeterPaul.nutrients$lakename:PeterPaul.nutrients$depth_id  6   1865502
## Residuals                        1409  95237896
##                                Mean Sq F value
## PeterPaul.nutrients$lakename      4034942    59.7
## PeterPaul.nutrients$depth_id      19281270   285.3
## PeterPaul.nutrients$lakename:PeterPaul.nutrients$depth_id  310917     4.6
## Residuals                        67593
##                                Pr(>F)
## PeterPaul.nutrients$lakename      2.09e-14 ***
## PeterPaul.nutrients$depth_id      < 2e-16 ***
## PeterPaul.nutrients$lakename:PeterPaul.nutrients$depth_id  0.000123 ***
```

```
## Residuals
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 922 observations deleted due to missingness
```

If the interaction is significant, we interpret pairwise differences for the interaction. If the interaction is not significant, we interpret differences for the main effects only.

```
TukeyHSD(TNanova.interaction)
```

```
## Tukey multiple comparisons of means
## 95% family-wise confidence level
##
## Fit: aov(formula = PeterPaul.nutrients$tn_ug ~ PeterPaul.nutrients$lakename * PeterPaul.nutrients$depth_id)
##
## $`PeterPaul.nutrients$lakename`
##              diff          lwr          upr p adj
## Peter Lake-Paul Lake 106.4994 79.45986 133.5389      0
##
## $`PeterPaul.nutrients$depth_id`
##              diff          lwr          upr          p adj
## 2-1  97.28178    22.187578 172.375977 0.0026048
## 3-1 113.40580    38.561123 188.250474 0.0001681
## 4-1  78.98288     6.027266 151.938490 0.0239560
## 5-1  22.46140   -54.472262  99.395056 0.9779694
## 6-1  39.00303   -47.439981 125.446032 0.8368008
## 7-1 859.47649  796.403376 922.549613 0.0000000
## 3-2  16.12402   -80.781877 113.029920 0.9989677
## 4-2 -18.29890 -113.753334  77.155534 0.9977033
## 5-2 -74.82038 -173.348628  23.707867 0.2736275
## 6-2 -58.27875 -164.398595  47.841091 0.6684093
## 7-2 762.19472  674.062737 850.326698 0.0000000
## 4-3 -34.42292 -129.681178  60.835337 0.9376228
## 5-3 -90.94440 -189.282605   7.393801 0.0914950
## 6-3 -74.40277 -180.346191  31.540645 0.3690493
## 7-3 746.07070  658.151229 833.990163 0.0000000
## 5-4 -56.52148 -153.429674  40.386713 0.6012407
## 6-4 -39.97985 -144.597267  64.637563 0.9194461
## 7-4 780.49362  694.176594 866.810640 0.0000000
## 6-5  16.54163  -90.887744 123.971001 0.9993372
## 7-5 837.01510  747.310610 926.719585 0.0000000
## 7-6 820.47347  722.491325 918.455613 0.0000000
##
## $`PeterPaul.nutrients$lakename:PeterPaul.nutrients$depth_id`
##              diff          lwr          upr          p adj
## Peter Lake:1-Paul Lake:1  143.1294811   73.915904 212.343059 0.0000000
## Paul Lake:2-Paul Lake:1   89.7855459  -30.633879 210.204971 0.4051921
## Peter Lake:2-Paul Lake:1  248.2658331  127.038156 369.493511 0.0000000
## Paul Lake:3-Paul Lake:1  101.4661113  -18.165112 221.097334 0.2014271
## Peter Lake:3-Paul Lake:1  269.3924944  148.164817 390.620172 0.0000000
## Paul Lake:4-Paul Lake:1   89.6427657  -27.021048 206.306579 0.3532527
## Peter Lake:4-Paul Lake:1  211.6260697   93.514201 329.737938 0.0000002
## Paul Lake:5-Paul Lake:1   56.1537708  -68.524801 180.832343 0.9653743
## Peter Lake:5-Paul Lake:1  132.3545363   9.446828 255.262245 0.0213490
## Paul Lake:6-Paul Lake:1  158.0133808  18.928611 297.098151 0.0104293
```

## Peter Lake:6-Paul Lake:1	63.0646474	-76.020123	202.149418	0.9634431
## Paul Lake:7-Paul Lake:1	928.9721994	827.295000	1030.649398	0.0000000
## Peter Lake:7-Paul Lake:1	933.5806988	832.307171	1034.854227	0.0000000
## Paul Lake:2-Peter Lake:1	-53.3439353	-173.794582	67.106712	0.9698240
## Peter Lake:2-Peter Lake:1	105.1363520	-16.122339	226.395043	0.1732123
## Paul Lake:3-Peter Lake:1	-41.6633698	-161.326020	77.999281	0.9966269
## Peter Lake:3-Peter Lake:1	126.2630133	5.004322	247.521705	0.0320229
## Paul Lake:4-Peter Lake:1	-53.4867154	-170.182756	63.209325	0.9601487
## Peter Lake:4-Peter Lake:1	68.4965885	-49.647112	186.640289	0.7995797
## Paul Lake:5-Peter Lake:1	-86.9757103	-211.684438	37.733018	0.5229580
## Peter Lake:5-Peter Lake:1	-10.7749448	-133.713243	112.163354	1.0000000
## Paul Lake:6-Peter Lake:1	14.8838996	-124.227903	153.995703	1.0000000
## Peter Lake:6-Peter Lake:1	-80.0648337	-219.176637	59.046969	0.8078319
## Paul Lake:7-Peter Lake:1	785.8427183	684.128544	887.556893	0.0000000
## Peter Lake:7-Peter Lake:1	790.4512176	689.140567	891.761868	0.0000000
## Peter Lake:2-Paul Lake:2	158.4802873	2.230523	314.730052	0.0429890
## Paul Lake:3-Paul Lake:2	11.6805655	-143.333848	166.694979	1.0000000
## Peter Lake:3-Paul Lake:2	179.6069485	23.357184	335.856713	0.0088082
## Paul Lake:4-Paul Lake:2	-0.1427801	-152.878776	152.593216	1.0000000
## Peter Lake:4-Paul Lake:2	121.8405238	-32.004374	275.685421	0.3031203
## Paul Lake:5-Paul Lake:2	-33.6317750	-192.573857	125.310307	0.9999857
## Peter Lake:5-Paul Lake:2	42.5689905	-114.987805	200.125786	0.9997652
## Paul Lake:6-Paul Lake:2	68.2278349	-102.249179	238.704849	0.9874061
## Peter Lake:6-Paul Lake:2	-26.7208984	-197.197912	143.756116	0.9999996
## Paul Lake:7-Paul Lake:2	839.1866536	697.567122	980.806185	0.0000000
## Peter Lake:7-Paul Lake:2	843.7951529	702.465161	985.125145	0.0000000
## Paul Lake:3-Peter Lake:2	-146.7997218	-302.442841	8.843397	0.0882601
## Peter Lake:3-Peter Lake:2	21.1266613	-135.746857	178.000180	0.9999999
## Paul Lake:4-Peter Lake:2	-158.6230674	-311.997108	-5.249027	0.0346196
## Peter Lake:4-Peter Lake:2	-36.6397634	-191.118126	117.838599	0.9999459
## Paul Lake:5-Peter Lake:2	-192.1120623	-351.667373	-32.556751	0.0043137
## Peter Lake:5-Peter Lake:2	-115.9112968	-274.086692	42.264099	0.4356459
## Paul Lake:6-Peter Lake:2	-90.2524523	-261.301347	80.796442	0.8881613
## Peter Lake:6-Peter Lake:2	-185.2011857	-356.250080	-14.152291	0.0199298
## Paul Lake:7-Peter Lake:2	680.7063663	538.398939	823.013793	0.0000000
## Peter Lake:7-Peter Lake:2	685.3148656	543.295577	827.334155	0.0000000
## Peter Lake:3-Paul Lake:3	167.9263831	12.283264	323.569502	0.0208421
## Paul Lake:4-Paul Lake:3	-11.8233456	-163.938683	140.291992	1.0000000
## Peter Lake:4-Paul Lake:3	110.1599583	-43.068773	263.388689	0.4694542
## Paul Lake:5-Paul Lake:3	-45.3123405	-203.658092	113.033411	0.9995597
## Peter Lake:5-Paul Lake:3	30.8884250	-126.066777	187.843627	0.9999939
## Paul Lake:6-Paul Lake:3	56.5472694	-113.373900	226.468439	0.9978542
## Peter Lake:6-Paul Lake:3	-38.4014639	-208.322634	131.519706	0.9999691
## Paul Lake:7-Paul Lake:3	827.5060881	686.556156	968.456020	0.0000000
## Peter Lake:7-Paul Lake:3	832.1145874	691.455574	972.773601	0.0000000
## Paul Lake:4-Peter Lake:3	-179.7497287	-333.123769	-26.375688	0.0065958
## Peter Lake:4-Peter Lake:3	-57.7664247	-212.244787	96.711937	0.9932710
## Paul Lake:5-Peter Lake:3	-213.2387236	-372.794035	-53.683412	0.0006485
## Peter Lake:5-Peter Lake:3	-137.0379581	-295.213354	21.137437	0.1741687
## Paul Lake:6-Peter Lake:3	-111.3791136	-282.428008	59.669781	0.6387944
## Peter Lake:6-Peter Lake:3	-206.3278470	-377.376741	-35.278953	0.0041875
## Paul Lake:7-Peter Lake:3	659.5797050	517.272278	801.887132	0.0000000
## Peter Lake:7-Peter Lake:3	664.1882044	522.168915	806.207493	0.0000000
## Peter Lake:4-Paul Lake:4	121.9833039	-28.940054	272.906662	0.2709820


```
## Paul Lake:5-Paul Lake:4      -33.4889949 -189.604955  122.626965 0.9999832
## Peter Lake:5-Paul Lake:4      42.7117706 -111.993599  197.417140 0.9997021
## Paul Lake:6-Paul Lake:4       68.3706150  -99.474611  236.215841 0.9852540
## Peter Lake:6-Paul Lake:4     -26.5781183 -194.423344  141.267107 0.9999996
## Paul Lake:7-Paul Lake:4      839.3294337  700.889196  977.769671 0.0000000
## Peter Lake:7-Paul Lake:4      843.9379330  705.793899  982.081967 0.0000000
## Paul Lake:5-Peter Lake:4    -155.4722989 -312.673320    1.728722 0.0560457
## Peter Lake:5-Peter Lake:4    -79.2715333 -235.071788    76.528721 0.9128009
## Paul Lake:6-Peter Lake:4    -53.6126889 -222.467620  115.242242 0.9986743
## Peter Lake:6-Peter Lake:4  -148.5614222 -317.416354    20.293509 0.1558154
## Paul Lake:7-Peter Lake:4     717.3461298  577.683438  857.008821 0.0000000
## Peter Lake:7-Peter Lake:4     721.9546291  582.585543  861.323715 0.0000000
## Peter Lake:5-Paul Lake:5       76.2007655  -84.634717  237.036248 0.9483731
## Paul Lake:6-Paul Lake:5     101.8596100  -71.652121  275.371341 0.7849894
## Peter Lake:6-Paul Lake:5       6.9108766 -166.600854  180.422608 1.0000000
## Paul Lake:7-Paul Lake:5     872.8184286  727.560037 1018.076820 0.0000000
## Peter Lake:7-Paul Lake:5     877.4269279  732.450809 1022.403047 0.0000000
## Paul Lake:6-Peter Lake:5      25.6588444 -146.584818  197.902507 0.9999998
## Peter Lake:6-Peter Lake:5    -69.2898889 -241.533551  102.953773 0.9868114
## Paul Lake:7-Peter Lake:5     796.6176631  652.876372  940.358954 0.0000000
## Peter Lake:7-Peter Lake:5     801.2261624  657.770129  944.682196 0.0000000
## Peter Lake:6-Paul Lake:6     -94.9487333 -279.084954    89.187487 0.9043108
## Paul Lake:7-Paul Lake:6     770.9588186  613.162028  928.755610 0.0000000
## Peter Lake:7-Paul Lake:6     775.5673180  618.030332  933.104304 0.0000000
## Paul Lake:7-Peter Lake:6     865.9075520  708.110761 1023.704343 0.0000000
## Peter Lake:7-Peter Lake:6     870.5160513  712.979065 1028.053037 0.0000000
## Peter Lake:7-Paul Lake:7       4.6084993 -121.135612  130.352610 1.0000000
```

Pairs are considered to be in the same grouping if the p-value for that pairing is > 0.05 . It is easy to see that this grouping process can become complicated when many factors are present for each variable! For a challenge, try writing code that will generate groupings for each factor level in the dataset using the `glht` function in the `multcomp` package.

Exercise

Run the same tests and visualizations (main and interaction effects two-way ANOVA) for total phosphorus concentrations. How do your results compare for the different nutrients?